



Rational Proofs with Multiple Provers^{*}

Jing Chen

Samuel McCauley

Shikha Singh

Computer Science Department, Stony Brook University
Stony Brook, NY 11794, USA
{jingchen, smccauley, shiksingh}@cs.stonybrook.edu

ABSTRACT

Interactive proofs model a world where a verifier delegates computation to an untrustworthy prover, verifying the prover's claims before accepting them. These proofs have applications to delegation of computation, probabilistically checkable proofs, crowdsourcing, and more.

In some of these applications, the verifier may pay the prover based on the quality of his work. Rational proofs, introduced by Azar and Micali (2012), are an interactive proof model in which the prover is *rational* rather than untrustworthy—he may lie, but only to increase his payment. This allows the verifier to leverage the greed of the prover to obtain better protocols: while rational proofs are no more powerful than interactive proofs, the protocols are simpler and more efficient. Azar and Micali posed as an open problem whether multiple provers are more powerful than one for rational proofs.

We provide a model that extends rational proofs to allow multiple provers. In this model, a verifier can cross-check the answers received by asking several provers. The verifier can pay the provers according to the quality of their work, incentivizing them to provide correct information.

We analyze rational proofs with multiple provers from a complexity-theoretic point of view. We fully characterize this model by giving tight upper and lower bounds on its power. On the way, we resolve Azar and Micali's open problem in the affirmative, showing that multiple rational provers are strictly more powerful than one (under standard complexity-theoretic assumptions). We further show that the full power of rational proofs with multiple provers can be achieved using only two provers and five rounds of interaction. Finally, we consider more demanding models where the verifier wants the provers' payment to decrease signifi-

cantly when they are lying, and fully characterize the power of the model when the payment gap must be noticeable (i.e., at least $1/p$ where p is a polynomial).

Categories and Subject Descriptors

F.1.2 [Computation by Abstract Devices]: Modes of Computation—*Interactive computation*; F.1.3 [Computation by Abstract Devices]: Complexity Measures and Classes; F.2.2 [Analysis of Algorithms and Problem Complexity]: Nonnumerical Algorithms and Problems—*Complexity of proof procedures*

Keywords

Interactive proofs; multi-prover rational interactive proofs; scoring rules; DC uniform circuit families; complexity theory

1. INTRODUCTION

Multi-prover interactive proofs (MIP) [10] and rational interactive proofs (RIP) [5] are two important extensions of interactive proof systems.

In a multi-prover interactive proof, several computationally unbounded, potentially dishonest provers interact with a polynomial time, randomized verifier. The provers can pre-agree on a joint strategy to convince the verifier of the truth of a proposition. However, once the protocol starts, the provers cannot communicate with each other. If the proposition is true, the verifier should be convinced with probability 1; otherwise the verifier should reject with some non-negligible probability. As shown by Babai, Fortnow and Lund, $MIP = NEXP$ [8], which demonstrates the power of multiple provers compared to one-prover interactive proofs (recall that $IP = PSPACE$ [31, 29]).

Rational interactive proofs [5] are a variant of interactive proofs, where the verifier makes a payment to the prover at the end of the protocol. The prover is assumed to be *rational*, that is, he only acts in ways that maximize this payment. Thus, in contrast to interactive proofs, the prover does not care whether the verifier is convinced or not. Rational proofs ensure that the prover's payment is maximized if and only if the verifier learns the correct answer to the proposition. Azar and Micali [5] show that, while rational proofs are no more powerful than interactive proofs (i.e., $RIP = PSPACE$), the protocols are simpler and more efficient. Previous work on rational proofs [6, 5, 25] considers a *single* rational prover.

Many computation-outsourcing applications have ingredients of both of these models: the verifier pays a *team*

^{*}We thank several anonymous reviewers for their valuable feedback, and Sanjoy Das, Andrew Drucker, Silvio Micali and Rafael Pass for their comments. The second and third authors are partially supported by NSF Grants CNS 1408695, CCF 1439084, IIS 1247726, IIS 1251137, and CCF 1217708, and Sandia National Laboratories.

Permission to make digital or hard copies of all or part of this work for personal or classroom use is granted without fee provided that copies are not made or distributed for profit or commercial advantage and that copies bear this notice and the full citation on the first page. Copyrights for components of this work owned by others than the author(s) must be honored. Abstracting with credit is permitted. To copy otherwise, or republish, to post on servers or to redistribute to lists, requires prior specific permission and/or a fee. Request permissions from Permissions@acm.org.

ITCS'16, January 14–16, 2016, Cambridge, MA, USA.

Copyright is held by the owner/author(s). Publication rights licensed to ACM.

ACM 978-1-4503-4057-1/16/01 ...\$15.00.

DOI: <http://dx.doi.org/10.1145/2840728.2840744>.

of provers based on their responses. For example, in Internet marketplaces such as *Amazon's Mechanical Turk* [1] and *Proof Market* [3], the requesters (verifiers) post labor-intensive tasks on the website along with a monetary compensation they are willing to pay. The providers (provers) accept these offers and perform the job. In these Internet marketplaces and crowdsourcing games [33] correctness is often ensured by verifying one provider's answers against another [32, 2]. Thus, the providers collaborate as a team—their answers need to match, even though they are likely to not know each other and cannot communicate with each other [27].

Inspired by these applications and previous theoretical work, we introduce *multi-prover rational interactive proofs*, which combine elements of rational proofs and classical multi-prover interactive proofs. This model aims to answer the following question: what problems can be solved by a team of rational workers who cannot communicate with each other and get paid based on the joint-correctness of their answers? One of our main contributions is to completely characterize the power of this model.

Previous Work Involving Multiple Rational Provers.

The notion of rational proofs with multiple provers has appeared several times in previous work [5, 6, 25]. However, the authors only use multiple provers to simplify the analysis of single-prover protocols, without formalizing the model. They show that multiple provers in their protocols can be simulated by a single prover by scaling the payments appropriately. Azar and Micali [5] discuss one of the fundamental challenges of using multiple rational provers: in a cooperative setting, one prover may lie to give subsequent provers the opportunity to obtain a larger payment. They pose the following open problem: are multiple provers more powerful than one in rational proofs? In this paper, we show that, in general, a protocol with multiple rational provers cannot be simulated by a single-prover protocol under standard complexity-theoretic assumptions.

The Model of Multi-Prover Rational Proofs.

We briefly summarize our model in order to compare it with the literature and discuss our results. The model is formally defined in Section 2.

In a multi-prover rational interactive proof (MRIP), several computationally-unbounded provers communicate with a polynomial-time randomized verifier who wants to determine the membership of an input string in a language. The provers can pre-agree on how they plan to respond to the verifier's queries. However, they cannot communicate with each other once the protocol begins. At the end of the protocol, the verifier outputs the answer and computes a total payment for the provers based on the input, his own randomness, and the messages exchanged. This total payment may be distributed in any pre-determined way by the verifier or the provers themselves.

A protocol is an MRIP protocol if any strategy of the provers that maximizes their expected payment leads the verifier to the correct answer. The class of languages having such protocols is denoted by MRIP.

Distribution of Payments.

In classical MIP protocols, the provers work collaboratively to convince the verifier of the truth of a proposition

and their goal is to maximize the verifier's acceptance probability. Similarly, the rational provers in MRIP work as a team to maximize the total payment received from the verifier. Any pre-specified distribution of this payment is allowed, as long as it does not depend on the transcript of the protocol (i.e., the messages exchanged, the coins flipped, and the amount of the payment). For instance, the division of the payment can be pre-determined by the provers themselves based on the amount of work each prover must perform, or it can be pre-determined by the verifier based on the reputation of each prover.¹ We ignore the choice of division in our model and protocols, as it does not affect the choices made by the provers in deciding their strategy.

1.1 Results and Discussions

We state our main results and discuss several interesting aspects of our model.

The Power of Multi-Prover Rational Proofs.

We give tight upper and lower bounds on the power of MRIP protocols. As a warm up, we show that MRIP contains NEXP, using an MIP protocol as a black box and paying the provers appropriately. A similar technique shows that MRIP also contains coNEXP. In fact, an important property of MRIP is that it is closed under complement (Lemma 1). Thus, MRIP is strictly more powerful than RIP (assuming $\text{PSPACE} \neq \text{NEXP}$), resolving the question raised in [5]. Furthermore, MRIP is also more powerful than MIP (assuming $\text{NEXP} \neq \text{coNEXP}$), in contrast to the single-prover case in which classical and rational proofs have the same power.




Figure 1: The relative power of rational and classical interactive proof systems. Note that it is widely believed that $\text{PSPACE} \neq \text{EXP}$, $\text{EXP} \neq \text{NEXP}$, and $\text{NEXP} \neq \text{coNEXP}$.

We exactly characterize the class MRIP by showing that a language has a multi-prover rational interactive proof if and only if it is decidable by a deterministic exponential-time oracle Turing machine with non-adaptive access to an NP oracle. That is,

THEOREM 1. $\text{MRIP} = \text{EXP}^{\text{||NP}}$.

¹Note that unbalanced divisions are allowed: for example, it may be that one particular prover always receives half of the total payment, and the other provers split the remainder equally.

We summarize the power of various models of interactive proofs in Figure 1. To prove Theorem 1, we (a) provide a new characterization for EXP^{MIP} , and (b) decompose the circuit family for EXP^{MIP} obtained from this characterization into three stages, construct MRIP protocols for each stage and combine them together appropriately. A similar 3-stage decomposition was used in [6], but their technique results in an exponential blow-up in the number of messages when applied to multi-prover protocols.

It is known that any MIP protocol can be simulated using only two provers and one round of communication between the provers and the verifier [15]. Similarly, we show that only two provers and five rounds are sufficient to capture the full power of MRIP. That is, denoting by $\text{MRIP}(p, r)$ the class of languages that have p -prover rational proofs with r rounds, we have:

THEOREM 2. $\text{MRIP} = \text{MRIP}(2, 5)$.

Note that we count the number of rounds as the total number of interactions (provers' messages and verifier's queries are separate rounds), in contrast to the number of *pairs* of back-and-forth interaction of MIP [15]. We use this convention to simplify our technical discussion.² It remains open whether or not the number of rounds in our protocol can be further reduced (see Section 4 for further discussion).

Utility Gaps.

Rational proofs assume that the provers always act to maximize their payment. However, how much do they lose by lying? The notion of *utility gaps*, first introduced in [6], measures the loss in payment (or utility) incurred by a lying prover. Notice that a lying prover may (a) deviate (even slightly) from the truthful protocol but still lead the verifier to the correct answer or (b) deviate and mislead the verifier to an incorrect answer. The authors of [6] demanded their protocols to be robust against provers of type (a): that is, any deviation from the prescribed strategy that the verifier intends the provers to follow results in a significant decrease in the payment. This ideal requirement on utility gaps is quite strong, and even the protocol in [6] fails to satisfy it, as pointed out by Guo, Hubáček, Rosen and Vald [25]. In this paper, we consider multi-prover rational proofs robust against cheating provers of type (b): that is, the provers may respond to some messages of the verifier incorrectly and incur a small payment loss, but if the verifier learns the answer to the membership question of the input string incorrectly, the provers must suffer a significant loss in payment. Note that [25] also considers type (b) utility gap, but for single-prover protocols. Our utility gaps are formally defined in Section 5.

We show that requiring a noticeable utility gap results in protocols for a different, possibly smaller, complexity class. In particular, let $\text{poly}(n)$ -gap-MRIP be the class of languages that have MRIP protocols with $1/\alpha(n)$ utility gap for lying provers, where $\alpha(n)$ is a polynomial in n . We completely characterize this class as the class of languages decidable by a polynomial-time oracle Turing machine with non-adaptive access to an NEXP oracle. That is,

THEOREM 3. $\text{poly}(n)$ -gap-MRIP = P^{NEXP} .

²Using the convention of [15], our simulation of any MRIP protocol with two provers would require only three rounds.

Simple and Efficient MRIP protocol for NEXP.

Classical multi-prover interactive proofs for languages in NEXP rely on the MIP protocol for the NEXP-complete language Oracle-3SAT [8]. This MIP protocol is (a) complicated, involving techniques such as multilinearity test of functions and arithmetization, and (b) requires polynomial computation and polynomial rounds of interaction from the verifier. We construct a simple two prover, three round MRIP protocol for Oracle-3SAT using *scoring rules*.³ Our protocol is very efficient: the verifier performs linear computation along with constant number of basic arithmetic operations to calculate the reward.

Contrasting MRIP with Other Relevant Models.

Existing models, such as *refereed games*, MIP, and single-prover rational proofs all differ from MRIP in distinct and potentially interesting ways.

Refereed games [17] are interactive proof models consisting of two competing provers, who try to convince the verifier of the membership (or non-membership) of a given input string in a language. However, one of them is honest and the other is not. The model of refereed games reflects the strategic nature of the provers, but does not allow collaboration between them.

Classical multi-prover interactive proofs are robust against arbitrary collaborative provers, who may be irrational or even malicious. While MIP protocols provide a stronger guarantee, they are often complicated and computationally-intensive (require polynomial-work from the verifier). In contrast, MRIP restricts provers to be rational, a reasonable assumption in a “mercantile world”, as pointed out in [5]. This restriction leads to simple and efficient protocols for the single-prover case [6, 25], making rational proofs a meaningful and interesting model to study.

Finally, MRIP achieves its full power with only five rounds of interaction between the verifier and the provers. In contrast, RIP is less powerful when restricted to constant rounds.⁴

1.2 Additional Related Work

Interactive Proofs.

First introduced by Goldwasser, Micali and Rackoff [23] and in a different form by Babai and Moran [7], interactive proofs (IP) have been extensively studied (see, e.g., [22, 8, 10, 9, 23, 19, 20, 24]) and precisely characterized: that is, $\text{IP} = \text{PSPACE}$ [31, 29]. Ben-Or et al. [10] introduced multi-prover interactive proofs (MIP), which were shown to be exactly equal to NEXP [8]. In fact, two provers and one round is sufficient; in other words, $\text{NEXP} = \text{MIP}(2, 1)$ [15].

Rational Proofs.

Azar and Micali [5] use scoring rules in a novel way to construct simple and efficient single-prover rational protocols. The prover is assumed to be sensitive to exponentially small losses in payment, that is, a negligible utility gap is a sufficient punishment for the prover. In [6], the same authors proposed the idea of utility gaps and constructed super-

³Scoring rules are powerful tools to elicit information about probabilistic distributions from experts; see Section 2.

⁴As shown in [5], $\text{RIP-O}(1) = \text{CH}$, where $\text{RIP-O}(1)$ is the class of languages having constant-round rational proofs and CH is the counting hierarchy.

efficient rational proofs, where the verifier performs only logarithmic computation. Guo, Hubáček, Rosen and Vald [25] studied *rational arguments*, which are rational proofs with a single computationally-bounded prover. They constructed rational arguments with single-round interaction and sub-linear verification for NC^1 . They also mentioned that a rational proof model with multiple provers may have interesting implications in computation delegation schemes, but did not define such a model. Later, they extended their rational argument protocols to P [26], via a novel technique called *rational sumcheck*.

Game-Theoretic Characterization of Complexity Classes.

Game-theoretic characterization of complexity classes has been largely studied in the form of *refereed games* [12, 17, 14, 16, 30, 18, 28].

The classic work of Chandra and Stockmeyer [12] proved that any language in PSPACE is refereeable by a game of perfect information. Feige and Kilian [14] show this is tight for single-round refereed games and that the class of languages with polynomial-round refereed games is exactly EXP .

Feigenbaum, Koller and Shor [18] study a related complexity class EXP^{NP} and show that it can be simulated as a zero-sum refereed game between two computationally unbounded provers with *imperfect recall*. Note that imperfect recall is a very strong assumption and makes the computationally unbounded provers essentially act as oracles. In contrast, MRIP consists of collaborative provers having imperfect information (since a prover does not see the messages exchanged between the verifier and other provers) and perfect recall (since a prover remembers the history of messages he exchanged with the verifier). Notice that imperfect information is necessary for multi-prover protocols: if all provers can see all messages exchanged in the protocol, then the model degenerates to a single-prover case. Moreover, perfect recall gives the provers the ability to cheat adaptively across messages. With these differences, MRIP is equivalent to EXP^{NP} . To our best knowledge, this is the first game-theoretic characterization of this complexity class.

It is worth pointing out that EXP^{NP} is also an important class in the study of circuit lower bounds [34]. It would be interesting to see if the related complexity class EXP^{NP} emerges in similar contexts.

1.3 Outline of the paper

The paper is organized as follows. In Section 2 we define the class of languages that have multi-prover rational interactive proofs, MRIP, and discuss some of its properties. We construct MRIP protocols for NEXP in the same section. We characterize the class MRIP in Section 3. In Section 4 we show how to simulate any MRIP protocol with two provers and five rounds. Finally, in Section 5 we define the notion of *utility gap* and characterize the power of MRIP protocols with a polynomial utility gap.

2. MULTI-PROVER RATIONAL INTERACTIVE PROOFS

In this section we define the model of multi-prover rational interactive proofs and demonstrate several important properties of the class of languages recognized by these proofs.

First, we describe how the verifier and the provers interact, and introduce necessary notations on the way. Let L be a language, x a string whose membership in L is to be decided, and $n = |x|$. An interactive protocol is a pair (V, \vec{P}) , where V is the *verifier* and $\vec{P} = (P_1, \dots, P_{t(n)})$ is the vector of *provers*, with $t(n)$ a polynomial.⁵ The verifier runs in polynomial time and flips private coins, whereas each P_i is computationally unbounded. The verifier and provers all know x . The verifier can communicate with each prover privately, but no two provers can communicate with each other. In a *round*, either each prover sends a message to the verifier, or the verifier sends a message to each prover, and these two cases alternate. Without loss of generality we assume the first round of messages are sent by the provers, and the first bit sent by P_1 , denoted by c , indicates whether $x \in L$ (corresponding to $c = 1$) or not ($c = 0$). Notice that c does not depend on the randomness used by V .

The length of each message and the number of rounds are polynomial in n . Let $p(n)$ be the number of rounds and r the random string used by V . For each $j \in \{1, 2, \dots, p(n)\}$, let m_{ij} be the message exchanged between V and P_i in round j . In particular, the first bit of m_{11} is c . The transcript that each prover P_i has seen at the beginning of each round j is $(m_{i1}, m_{i2}, \dots, m_{i(j-1)})$. Let \vec{m} be the vector of all messages exchanged in the protocol (therefore \vec{m} is a random variable depending on r).

At the end of the communication, the verifier computes a payment function R based on x , r , and \vec{m} as the total payment to give to the provers. We restrict $R(x, r, \vec{m}) \in [-1, 1]$ for convenience.⁶

The protocol followed by V , including the payment function R , is public knowledge.

The verifier outputs c as the answer for the membership of x in L : that is, V does not check the provers' answer and follows it blindly. As will become clear from our results, this requirement for the verifier does not change the set of languages that have multi-prover rational interactive proofs. Indeed, we could have allowed V to compute his answer based on x , r , and \vec{m} as well, but the current model eases later discussion on the payment loss of the provers caused by "reporting a wrong answer".

Each prover P_i can choose a *strategy* $s_{ij} : \{0, 1\}^* \rightarrow \{0, 1\}^*$ for each round j , which maps the transcript he has seen at the beginning of round j to the message he sends in that round.⁷ Let $\vec{s}_i = (s_{i1}, \dots, s_{ip(n)})$ be the vector of strategies P_i uses in rounds $1, \dots, p(n)$, and $\vec{s} = (\vec{s}_1, \dots, \vec{s}_{t(n)})$ be the strategy profile of the provers. Given any input x , randomness r , and strategy profile \vec{s} , we denote by $(V, \vec{P})(x, r, \vec{s})$ the vector of all messages exchanged in the protocol.

The provers are *cooperative* and jointly act to maximize the (total) expected payment received from the verifier. Note that this is equivalent to the provers maximizing their own expected payment when each prover receives a pre-specified division of the payment, that is, for any function γ_i fixed at the beginning of the protocol with $\sum_{i=1}^n \gamma_i = 1$, P_i receives $\gamma_i R$.

⁵That is, we allow polynomially many provers.

⁶Note that the payment can be shifted and scaled so that it is in $[0, 1]$. We use both positive and negative payments in our model to better reflect the intuition behind our protocols: the former are rewards while the latter are punishments.

⁷ P_i does not send any message s_{ij} for an even-numbered round j , and these s_{ij} 's can be treated as constant functions.

Thus, before the protocol starts, the provers pre-agree on a strategy profile \tilde{s} that maximizes

$$u_{(V, \vec{P})}(\tilde{s}; x) \triangleq \mathbb{E}_r R(x, r, (V, \vec{P})(x, r, \tilde{s})).$$

When (V, \vec{P}) and x are clear from the context, we write $u(\tilde{s})$ for $u_{(V, \vec{P})}(\tilde{s}; x)$. Using the above notions, we define MRIP as follows.

DEFINITION 1 (MRIP). *For any language L , an interactive protocol (V, \vec{P}) is a multi-prover rational interactive proof (MRIP) protocol for L if, for any $x \in \{0, 1\}^*$ and any strategy profile \tilde{s} of the provers such that $u(\tilde{s}) = \max_{\tilde{s}'} u(\tilde{s}')$, we have*

1. $u(\tilde{s}) \geq 0$;
2. $c = 1$ if and only if $x \in L$.

We denote the set of languages that have MRIP protocols by MRIP, and the set of languages that have MRIP protocols with k provers and p rounds by MRIP(k, p).

LEMMA 1. *MRIP is closed under complement.*

PROOF. Let L be a language in MRIP, (V, \vec{P}) an MRIP protocol for L , and R the payment function computed by V . We construct a verifier V' and thus an MRIP protocol (V', \vec{P}) for \bar{L} , as follows.

- V' runs V to compute the messages he should send in each round, except that, whenever V' gives V as an input the first message m'_{11} sent by P_1 , he flips the first bit.
- At the end of the communication, V' computes a payment function R' : for any x, r , and \vec{m}' , $R'(x, r, \vec{m}') = R(x, r, \vec{m})$, where \vec{m} is \vec{m}' with the first bit flipped.
- V' outputs the first bit sent by P_1 .

For each strategy profile \tilde{s} of the provers in the protocol (V, \vec{P}) , consider the following strategy profile \tilde{s}' in the protocol (V', \vec{P}) :

- $\tilde{s}'_i = \tilde{s}_i$ for each $i \neq 1$.
- In round 1, \tilde{s}'_1 outputs the same message as \tilde{s}_1 , except that the first bit is flipped.
- For any odd number $k > 1$ and any transcript m'_1 for P_1 at the beginning of round k , $\tilde{s}'_1(m'_1)$ is the same as $\tilde{s}_1(m_1)$, where m_1 is m'_1 with the first bit flipped.

It is easy to see that for any x and r , $(V', \vec{P})(x, r, \tilde{s}')$ is the same as $(V, \vec{P})(x, r, \tilde{s})$ except the first bit. Thus $R'(x, r, (V', \vec{P})(x, r, \tilde{s}')) = R(x, r, (V, \vec{P})(x, r, \tilde{s}))$, which implies

$$u_{(V', \vec{P})}(\tilde{s}'; x) = u_{(V, \vec{P})}(\tilde{s}; x).$$

Also, it is easy to see that the mapping from \tilde{s} to \tilde{s}' is a bijection. Accordingly, arbitrarily fixing a strategy profile \tilde{s}' that maximizes $u_{(V', \vec{P})}(\tilde{s}'; x)$, we have that the corresponding \tilde{s} maximizes $u_{(V, \vec{P})}(\tilde{s}; x)$ as well. Thus

$$u_{(V', \vec{P})}(\tilde{s}'; x) = u_{(V, \vec{P})}(\tilde{s}; x) \geq 0,$$

where the inequality is by Definition 1. Furthermore, $x \in L$ if and only if the first bit sent by \tilde{s}_1 is 1, if and only if the first bit sent by \tilde{s}'_1 is 0. That is, $x \in \bar{L}$ if and only if the first bit sent by \tilde{s}'_1 is 1. \square

To demonstrate the power of multi-prover rational proofs, we show that MRIP contains NEXP. With Lemma 1, this implies that MRIP contains coNEXP as well. We show this in two different ways.

First, we construct an MRIP protocol for any language in NEXP using a standard MIP protocol as a subroutine. Given existing MIP protocols, this method is intuitive and its correctness is easy to see. However, the computation and communication complexity of the protocol depends on the MIP protocol.

Second, we construct an MRIP protocol for an NEXP-complete language without relying on MIP protocols, instead by exploiting the rational nature of the provers. This protocol uses *proper scoring rules* to incentivize the provers. In contrast to MIP, this MRIP protocol is very efficient: it only requires the verifier to perform linear amount of computation and communication, along with computing the payment using constant number of arithmetic operations.

2.1 An MRIP Protocol for any Language in NEXP, Based on MIP.

An MIP protocol (see, e.g., [8, 15]) for a language $L \in \text{NEXP}$ first reduces L to the NEXP-complete problem Oracle-3SAT (defined below), and then runs an MIP protocol for Oracle-3SAT.

DEFINITION 2 (Oracle-3SAT [8]). *Let w be a binary string of length $r + 3s$. Let B be a 3-CNF of $r + 3s + 3$ variables. A Boolean function $A : \{0, 1\}^s \rightarrow \{0, 1\}$ is a 3-satisfying oracle for B if $B(w, A(b_1), A(b_2), A(b_3))$ is satisfied for all w , where $b_1 b_2 b_3$ are the last $3s$ bits of w . The Oracle-3SAT problem is to decide, for a given B , if there is a 3-satisfying oracle for B .*

Our MRIP protocol for NEXP uses a black-box MIP protocol and an appropriate payment scheme.

LEMMA 2. $\text{NEXP} \subseteq \text{MRIP}$.

PROOF. The MRIP protocol (V, \vec{P}) for L is shown in Figure 2. By construction, the payment to the provers is always non-negative.

Now we show that V outputs 1 if and only if $x \in L$. On the one hand, for any $x \in L$, if the provers send $c = 1$ and execute the MIP protocol with V , then their payment is $R = 1$;⁸ while if they send $c = 0$, the payment is $R = 1/2 < 1$. Accordingly, their best strategy profile is to send $c = 1$ and run the MIP protocol correctly. On the other hand, for any $x \notin L$, if the provers send $c = 1$ and run the MIP protocol, then by the definition of MIP, the probability that V accepts is at most $1/3$, and the expected payment is at most $1/3$; while if they send $c = 0$, the payment is $1/2 > 1/3$. Accordingly, their best strategy profile is to send $c = 0$. \square

REMARK 1. *Since any language $L \in \text{NEXP}$ has a 2-prover 1-round MIP protocol [15], we automatically obtain an MRIP protocol for L with 2 provers and 3 rounds of interaction (using our convention, executing the MIP protocol takes 2 rounds, plus one round for the answer bit c). Note that this is the best possible. However, the computation and communication complexity of the protocol are both polynomial (in addition to the reduction to Oracle-3SAT) and involve complex techniques such as arithmetization and multilinearity test [8].*

⁸If the MIP protocol does not have perfect completeness and accepts x with probability at least $2/3$, then the expected payment is at least $2/3$. However, this does not affect the correctness of the MRIP protocol.

For any input string x , the protocol (V, \vec{P}) works as follows:

1. P_1 sends a bit c to V . V outputs c at the end of the protocol.
2. If $c = 0$ then the protocol ends and the payment given to the provers is $R = 1/2$;
3. Otherwise, V and \vec{P} run an MIP protocol for proving $x \in L$. If the verifier accepts then $R = 1$, else $R = 0$.

Figure 2: A simple MRIP protocol for NEXP.

By Lemmas 1 and 2, we immediately have the following corollary. Also note that we can obtain a 2-prover 3-round MRIP protocol for coNEXP directly by modifying the protocol of Lemma 2.

COROLLARY 1. $\text{coNEXP} \subseteq \text{MRIP}$.

2.2 An MRIP Protocol for any Language in NEXP, Using Scoring Rules.

We now construct an MRIP protocol for any language in NEXP without relying on MIP protocols. Instead, we use *proper scoring rules* to compute the payment that should be given to the provers so as to ensure truthful answers.

To begin, we recall the definitions of proper scoring rules in general and the *Brier's scoring rule* in particular, which has been an essential ingredient in the construction of rational proofs [5, 6, 25].

Proper Scoring Rules.

Scoring rules are tools to assess the quality of a probabilistic forecast by assigning a numerical score (that is, a payment to the forecaster) to it based on the predicted distribution and the sample that materializes. More precisely, given any probability space Σ , letting $\Delta(\Sigma)$ be the set of probability distributions over Σ , a *scoring rule* is a function from $\Delta(\Sigma) \times \Sigma$ to \mathbb{R} , the set of reals. A scoring rule S is *proper* if, for any distribution D over Σ and distribution $D' \neq D$, we have

$$\sum_{\omega \in \Sigma} D(\omega) S(D, \omega) \geq \sum_{\omega \in \Sigma} D(\omega) S(D', \omega),$$

where $D(\omega)$ is the probability that ω is drawn from D . A scoring rule S is *strictly proper* if the above inequality is strict. Notice that, when the true distribution is D , the forecaster maximizes the expected payment under a strictly proper scoring rule by reporting $D' = D$. For a comprehensive survey on scoring rules, see [21].

Brier's Scoring Rule [11].

This classic scoring rule, denoted by BSR , is defined as follows: for any distribution D and $\omega \in \Sigma$,

$$\text{BSR}(D, \omega) = 2D(\omega) - \sum_{\omega \in \Sigma} D(\omega)^2 - 1.$$

It is well known that BSR is strictly proper.

Notice that BSR requires the computation of $\sum_{\omega \in \Sigma} D(\omega)^2$, which can be hard when $|\Sigma|$ is large. However, similar to [5, 25], in this paper we only consider $\Sigma = \{0, 1\}$. Also notice that, BSR has range $[-2, 0]$ and can be shifted and scaled so that (1) the range is non-negative and bounded, and (2) the resulting scoring rule is still strictly proper. In particular, we shall add 2 to the function when using it.

Next, we construct a simple and efficient MRIP protocol for Oracle-3SAT . Similar to the classical multi-prover case, an MRIP protocol for any language $L \in \text{NEXP}$ can be obtained by first reducing L to Oracle-3SAT , and then using our efficient MRIP protocol for Oracle-3SAT . The complexity of the overall protocol for L is thus the same as the reduction.

LEMMA 3. *Oracle-3SAT has an MRIP protocol with 2 provers and 3 rounds where, for any instance B of length n , the randomness used by the verifier, the computation complexity, and the communication complexity of the protocol are all $O(n)$, and the computation of the payment function consists of constant number of arithmetic operations over $O(n)$ -bit numbers.*

PROOF. For any instance B with $r+3s+3$ variables (thus $n \geq r+3s+3$), the provers can, with their unbounded computation power, find an oracle A^* that maximizes the number of satisfying w 's for B . Denote this number by a^* . If $B \in \text{Oracle-3SAT}$ then $a^* = 2^{r+3s}$, otherwise $a^* < 2^{r+3s}$.

Roughly speaking, in our MRIP protocol (V, \vec{P}) , the verifier incentivizes the provers to report the correct value of a^* , so that the membership of B can be decided. The protocol is shown in Figure 3.

To see why this protocol works, first notice that, if the provers send $c = 0$ and $a = 0$ in Step 1 and always send 0's in Step 4, then the protocol always ends in Steps 5b or 5c, and the provers' expected payment is 0. Accordingly, the best strategy profile \tilde{s}^* of the provers gives them expected payment at least 0, and Condition 1 of Definition 1 holds. Moreover, \tilde{s}^* must be such that

$$\text{either } c = 1 \text{ and } a = 2^{r+3s}, \text{ or } c = 0 \text{ and } a < 2^{r+3s}, \quad (1)$$

since otherwise the provers' expected payment is -1 .

It remains to show that Condition 2 of Definition 1 holds. Notice that P_2 only answers one query of the verifier (in step 4), thus under any strategy \tilde{s}_2 and given any c and a , P_2 de facto commits to an oracle $A' : \{0, 1\}^s \rightarrow \{0, 1\}$. Assume that P_1 , using a strategy \tilde{s}_1 and seeing (b_1, \dots, b_6) , sends V six bits in step 4 that are not consistent with A' —that is, there exists $i \in \{1, \dots, 6\}$ such that $A(b_i) \neq A'(b_i)$. Let q be the probability that, conditioned on (b_1, \dots, b_6) , the verifier chooses a k that catches the provers in step 5a. We have $q \geq 1/6$. Let R be the payment to the provers conditioned on (b_1, \dots, b_6) and on the event that they are not caught in step 5a. Note that $R \leq \frac{2}{11}$ by the definition of Brier's scoring rule (after shifting and scaling). Thus the expected payment to the provers conditioned on (b_1, \dots, b_6) is $-q + (1-q)R < 0$. However, if P_1 answers the verifier's queries consistently with A' , his expected payment conditioned on (b_1, \dots, b_6) is non-negative. Accordingly, the best strategy profile \tilde{s}^* must be such that, for any c , a , and the oracle committed by P_2 , P_1 's answers for any (b_1, \dots, b_6) are consistent with A' . Thus under \tilde{s}^* the payment is never computed in step 5a.

Furthermore, given b_1, b_2, b_3 and A' , whether B evaluates to 0 in step 5b or not is totally determined. If B evaluates

to 0, then it does not matter what a or c is and the provers' received payment is 0. If B does not evaluate to 0 in step 5b, then the expected payment to the provers in step 5c is defined by Brier's scoring rule: the true distribution of b , denoted by D , is such that $D(1) = a'/2^{r+3s}$, with a' being the number of satisfying w 's for B under oracle A' ; and the realized value is $b = B(z', b_4, b_5, b_6, A(b_4), A(b_5), A(b_6))$. Indeed, since b_4, b_5, b_6 are independent from b_1, b_2, b_3 , we have that w' is a uniformly random input to B , and the probability for b to be 1 is exactly $a'/2^{r+3s}$. Since Brier's scoring rule is strictly proper, conditioned on A' the provers maximize their expected utility by reporting

$$a = a', \quad (2)$$

which implies $(p_1, p_0) = (D(1), D(0))$.

If $B \notin \text{Oracle-3SAT}$, then no matter which oracle A' is committed under \tilde{s}^* , we have $a' < 2^{r+3s}$. By Equations 1 and 2, we have $a < 2^{r+3s}$ and $c = 0$, as desired.

If $B \in \text{Oracle-3SAT}$, we show that under \tilde{s}^* the prover P_2 commits to the desired 3-satisfying oracle A^* (so that $a' = 2^{r+3s}$ and $D(1) = 1$). To see why this is true, denote by $\text{BSR}(D)$ the expected score for reporting D under Brier's scoring rule, when the true distribution is D . We have

$$\begin{aligned} \text{BSR}(D) &= D(1) (2D(1) - D(1)^2 - (1 - D(1))^2 - 1) \\ &\quad + (1 - D(1)) (2(1 - D(1)) - D(1)^2 \\ &\quad - (1 - D(1))^2 - 1) \\ &= 2(D(1)^2 - D(1)). \end{aligned}$$

Thus $\text{BSR}(D)$ is symmetric at $D(1) = 1/2$, strictly decreasing on $D(1) \in [0, 1/2]$, strictly increasing on $D(1) \in [1/2, 1]$, and maximized when $D(1) = 1$ or $D(1) = 0$. Notice that the shifting and scaling of BSR in step 5c do not change these properties, but make $\text{BSR}(D)$ strictly positive when $D(1) = 1$ or $D(1) = 0$. Accordingly, to maximize their expected payment conditioned on the event that step 5c is reached, P_2 should commit to either an oracle A' such that $D(1)$ is as small as possible, or an A' such that $D(1)$ is as large as possible, whichever makes $D(1)$ further from $1/2$.

If there is no oracle A' such that $a' = 0$ given A' , then the only way for the provers to maximize their expected payment is to commit to the 3-satisfying oracle A^* (thus $a' = 1$), under which step 5c is reached with probability 1. Again by Equations 1 and 2, we have $c = 1$ and $a = 2^{r+3s}$ as desired.

If there is both a 3-satisfying oracle A^* and an oracle A' such that $a' = 0$, we need to make sure that P_2 does not commit to A' : indeed, committing to any other oracle results in an expected payment strictly smaller than that by committing to A^* , since it increases the probability that the protocol ends at step 5b with $R = 0$, and strictly decreases the expected payment conditioned on step 5c being reached. If P_2 commits to A' , then B always evaluates to 0 in step 5b, and step 5c is actually never reached. Thus, even though by committing to A' the provers maximize their expected payment in step 5c, their actual expected payment is 0. Instead, by committing to A^* , step 5c is reached with probability 1 and the provers get positive payment. Accordingly, the strategy profile \tilde{s}^* must be such that P_2 commits to A^* and P_1 sends $a = 2^{r+3s}$ and $c = 1$, as desired. Notice that if there are multiple 3-satisfying oracles for B , then the provers can pre-agree on any one of them.

In sum, Condition 2 of Definition 1 also holds and (V, \vec{P}) is an MRIP protocol for **Oracle-3SAT**. Moreover, since $n \geq r+3s+3$, it is easy to see that the number of coins flipped by V is at most $2n$ (for sampling w , w' , and k) and the number of bits exchanged between V and \vec{P} is at most $3n+3$. Finally, given an input string $w = (z, b_1, b_2, b_3)$ for B and the 3-bit answers of the oracle for b_1, b_2, b_3 , B can be evaluated in linear time, thus the running time of V is $O(n)$ plus constant number of arithmetic operations to compute the payment in step 5c. Therefore Lemma 3 holds. \square

3. CHARACTERIZING MULTI-PROVER RATIONAL INTERACTIVE PROOFS

In this section we prove Theorem 1, that is,

THEOREM 1. $\text{MRIP} = \text{EXP}^{\text{INP}}$.

We first show that **MRIP** is the same as another complexity class, $\text{EXP}^{\text{poly-NEXP}}$, which we define below. We complete the proof of Theorem 1 by showing $\text{EXP}^{\text{INP}} = \text{EXP}^{\text{poly-NEXP}}$.

DEFINITION 3. $\text{EXP}^{\text{poly-NEXP}}$ is the class of languages decidable by a deterministic exponential-time Turing machine with non-adaptive access to an **NEXP** oracle, such that the length of each oracle query is polynomial in the length of the input of the Turing machine.

We use two intermediate lemmas to prove the lower bound. Lemma 4 is from [4] and gives a circuit characterization of **EXP**.

LEMMA 4 ([4]). *For any language L , $L \in \text{EXP}$ if and only if it can be computed by a DC uniform circuit family of size $2^{n^{O(1)}}$.*

LEMMA 5. *Every language L in **EXP** has an MRIP protocol with two provers and five rounds.*

PROOF. To begin, we recall some definitions and results from the literature that we use in our proof. First of all, a *circuit family* $\{C_n\}_{n=1}^\infty$ is a sequence of boolean circuits such that $C_n : \{0, 1\}^n \rightarrow \{0, 1\}$. The gates in the circuits are of type AND, OR, and NOT, with fan-ins 2, 2, and 1 respectively. The input string to a circuit is connected to a special set of "input gates", one for each bit of the input, whose output value is always the value of the corresponding bit. The *size* of a circuit C is the number of gates (including the input gates) in C . For a circuit C of size g , the set of gates in it is denoted by $\{1, 2, \dots, g\}$. Without loss of generality we assume that gate g is the output gate of the whole circuit. Moreover, if C has input length n , without loss of generality we assume that gates $1, 2, \dots, n$ are the input gates. Notice that the number of wires in C is at most $2g$, since each gate has at most 2 fan-ins. The set of wires is denoted by $\{1, 2, \dots, 2g\}$.

DEFINITION 4 (DC UNIFORM CIRCUITS [4]). *A circuit family $\{C_n\}_{n=1}^\infty$ is a Direct Connect uniform (DC uniform) family if the following questions can be answered in polynomial time:*

1. $\text{SIZE}(n)$: what is the size of C_n ?
2. $\text{INPUT}(n, h, i)$: is wire h an input to gate i in C_n ?
3. $\text{OUTPUT}(n, h, i)$: is wire h the output of gate i in C_n ?

For any instance B , the protocol (V, \vec{P}) works as follows:

1. P_1 sends $c \in \{0, 1\}$ and $a \in \{0, \dots, 2^{r+3s}\}$ to V . V always outputs c at the end of the protocol.
2. If $c = 1$ and $a < 2^{r+3s}$, or if $c = 0$ and $a = 2^{r+3s}$, the protocol ends with payment $R = -1$.
3. Otherwise, V randomly chooses two binary strings of length $r + 3s$, $w = (z, b_1, b_2, b_3)$ and $w' = (z', b_4, b_5, b_6)$, as well as a number $k \in \{1, 2, 3, 4, 5, 6\}$.
 V sends $b_1, b_2, b_3, b_4, b_5, b_6$ to P_1 and b_k to P_2 .
4. P_1 sends to V six bits, $A(b_i)$ with $i \in \{1, 2, \dots, 6\}$, and P_2 sends one bit, $A'(b_k)$.
5. The protocol ends and V computes the payment R as follows.
 - (a) If $A(b_k) \neq A'(b_k)$ then $R = -1$.
 - (b) Otherwise, if $B(z, b_1, b_2, b_3, A(b_1), A(b_2), A(b_3)) = 0$ then $R = 0$.
 - (c) Else, let $b = B(z', b_4, b_5, b_6, A(b_4), A(b_5), A(b_6))$, $p_1 = a/2^{r+3s}$, and $p_0 = 1 - p_1$.
 V computes R using Brier's scoring rule:
if $b = 1$ then $R = \frac{2p_1 - (p_1^2 + p_0^2) + 1}{11}$; otherwise $R = \frac{2p_0 - (p_1^2 + p_0^2) + 1}{11}$.

Figure 3: An efficient MRIP protocol for Oracle-3SAT.

4. TYPE(n, i, t): is t the type of gate i in C_n ?

That is, the circuits in a DC uniform family may have exponential size, but they have a succinct representation in terms of a polynomial-time Turing machine that can answer all the questions in Definition 4. The class EXP can be characterized by the class of DC uniform circuit families [4]; this is stated formally in Lemma 4.

Now we are ready to prove Lemma 5.

Following Lemma 4, there exists a DC uniform circuit family $\{C_n\}_{n=1}^\infty$ that computes L . Let $g = 2^{n^k}$ be the size of each C_n , where k is a constant that may depend on L . For any input string x of length n and any gate $i \in \{1, 2, \dots, g\}$ in C_n , let $v_i(x) \in \{0, 1\}$ be the value of i 's output on input x . In particular, $v_i(x) = x_i$ for any $i \in \{1, 2, \dots, n\}$. We call a gate i' in C_n an *input gate* of i if there is a directed wire from i' to i . The MRIP protocol (V, \vec{P}) is shown in Figure 4.

Clearly this protocol has two provers and five rounds. To see why it is an MRIP protocol, notice that if P_1 and P_2 always send the correct c and answer V 's queries correctly according to C_n , their received payment is always $R = 1$, irrespective of V 's coin flips. Thus the expected received payment is 1. Below we show that any other strategy profile gives expected payment strictly less than 1.

First of all, when the gate i chosen by the verifier in step 2 is not an input gate, if any of P_1 's answers in step 3 to queries 2a and 2b (namely, about i 's type, input gates and input wires) is incorrect, then the verification in step 6a fails, giving the provers a payment $R = 0$. Accordingly, if there exists such a gate then the expected payment to the provers is at most $1 - 1/g < 1$. Similarly, if there exists a non-input gate i such that P_1 answers queries 2a and 2b correctly but the values $v_i(x), v_{i_1}(x), v_{i_2}(x)$ are inconsistent with i 's type, then, when i is chosen, step 6b fails, thus the expected payment to the provers is at most $1 - 1/g < 1$. Moreover, if there exists an input gate i such that $v_i(x) \neq x_i$ when i is chosen, or if $v_g(x) \neq c$ when gate g is chosen, then the expected payment to the provers is again at most $1 - 1/g < 1$.

Next, similar to the analysis of Lemma 3, P_2 is only queried once (in step 5). Thus P_2 de facto commits to an oracle $A : \{1, \dots, g\} \rightarrow \{0, 1\}$, which maps any gate to its value under input x . If there exists a gate i such that the values $v_i(x), v_{i_1}(x), v_{i_2}(x)$ in step 3 are not consistent with A , then, conditioned on i being chosen in step 2, with prob-

ability $1/3$ step 6e fails. Since i is chosen with probability $1/g$, the expected payment is at most $1 - \frac{1}{3g} < 1$.

Accordingly, the only strategy profile \vec{s} that can have expected payment equal to 1 is exactly what we want: that is,

- (1) P_1 and P_2 answer queries to the values of gates using the same oracle $A : \{1, \dots, g\} \rightarrow \{0, 1\}$,
- (2) $A(i) = x_i$ for any input gate i ,
- (3) $A(g) = c$, and
- (4) for any non-input gate i , $A(i)$ is computed correctly according to i 's type and the values of its input gates in C_n .

Thus $A(g)$ is computed according to C_n with input x , and $A(g) = 1$ if and only if $x \in L$. Since $c = A(g)$, we have that $c = 1$ if and only if $x \in L$, implying that (V, \vec{P}) is an MRIP protocol for L . Therefore Lemma 5 holds. \square

Using these lemmas, Lemma 6 proves the lower bound on MRIP. We describe the main ideas in the proof sketch below. A detailed proof can be found in the full version of this paper [13].

LEMMA 6. $\text{EXP}^{\|\text{poly}\|-\text{NEXP}} \subseteq \text{MRIP}$.

PROOF SKETCH: Using the characterization of EXP in terms of DC uniform circuits from Lemma 4, we have an MRIP protocol for EXP with 2 provers and 5 rounds as in Lemma 5 (see Figure 4). We then combine this protocol with the MRIP protocol for Oracle-3SAT in Figure 2 to obtain the desired MRIP protocol for $\text{EXP}^{\|\text{poly}\|-\text{NEXP}}$.⁹ In particular, we use the protocol for Oracle-3SAT to answer NEXP oracle queries. A similar structure has been used by [6] for single-prover rational proofs. However, the prover in [6] can send the entire proof of a circuit satisfiability problem to the verifier, which is not feasible for us because our circuit may be exponentially large while our verifier is still of polynomial time. We overcome this problem by using a second prover to cross-check whether the first prover is answering questions about the circuit correctly. \square

Lemma 7 then shows that the above lower bound is tight, leading to an exact characterization.

⁹The MRIP protocol for Oracle-3SAT in Figure 3 can also be used, with appropriate modifications in the computation of the payment. We use the protocol in Figure 2 to simplify the analysis.

Given any string x of length n ,

1. P_1 sends one bit $c \in \{0, 1\}$ to V . V always outputs c at the end of the protocol.
 2. V computes $g = \text{SIZE}(n)$, picks a gate $i \in \{1, 2, \dots, g\}$ uniformly at random, and sends i to P_1 . That is, V queries P_1 for:
 - (a) the type of gate i ,
 - (b) the input gates of i and corresponding input wires, and
 - (c) the values of gate i and its input gates.
 3. P_1 sends to V the concatenation of the following strings: type $t_i \in \{AND, OR, NOT\}$; gates $i_1, i_2 \in \{1, 2, \dots, g\}$; wires $h_1, h_2 \in \{1, 2, \dots, 2g\}$; and values $v_i(x), v_{i_1}(x), v_{i_2}(x) \in \{0, 1\}$.
 4. V picks a gate $i' \in \{i, i_1, i_2\}$ uniformly at random and sends i' to P_2 .
 5. P_2 sends $v_{i'}(x) \in \{0, 1\}$ to V .
 6. The protocol ends and V computes the payment R by verifying the following statements:
 - (a) If $i \notin \{1, 2, \dots, n\}$ (that is, i is not an input gate), then $\text{TYPE}(n, i, t_i) = 1$, $\text{INPUT}(n, h_1, i) = \text{OUTPUT}(n, h_1, i_1) = 1$, and, if $t_i \neq NOT$, $\text{INPUT}(n, h_2, i) = \text{OUTPUT}(n, h_2, i_2) = 1$.
 - (b) If $i \notin \{1, 2, \dots, n\}$, then $v_i(x)$ is computed correctly based on type t_i from $v_{i_1}(x)$ and (when $t_i \neq NOT$) $v_{i_2}(x)$.
 - (c) If $i \in \{1, 2, \dots, n\}$ then $v_i(x) = x_i$.
 - (d) If $i = g$ (that is, the output gate of the circuit) then $v_i(x) = c$.
 - (e) The answers of P_1 and P_2 on the value of gate i' are consistent.
- If any of these verifications fails then $R = 0$; otherwise $R = 1$.

Figure 4: An MRIP protocol for EXP.

LEMMA 7. $\text{MRIP} \subseteq \text{EXP}^{\text{poly}}\text{-NEXP}$.

PROOF. Arbitrarily fix a language $L \in \text{MRIP}$ and let (V, \vec{P}) be an MRIP protocol for L . Since V runs in polynomial time, there exists a constant k such that, for any two payments R and R' that can be output by V under some input of length n and some randomness, we have

$$R \neq R' \Rightarrow |R - R'| \geq \frac{1}{2^{n^k}}.$$

For example, n^k can be an upper bound of V 's running time. Moreover, since V uses polynomially many random coins, there exists another constant k' such that, when a payment appears with positive probability under some input of length n , it must appear with probability at least $\frac{1}{2^{n^{k'}}$. Accordingly, for any input x of length n and any two strategy profiles \vec{s} and \vec{s}' of the provers, if the expected payments $u(\vec{s}; x)$ and $u(\vec{s}'; x)$ are different, then

$$|u(\vec{s}; x) - u(\vec{s}'; x)| \geq \frac{1}{2^{n^{k+k'}}}. \quad (3)$$

Consider the following deterministic oracle Turing machine M : Given any input x of length n , it divides the interval $[0, 1]$ to $2 \cdot 2^{n^{k+k'}}$ sub-intervals of length $\frac{1}{2 \cdot 2^{n^{k+k'}}}$.

For any $i \in \{1, \dots, 2 \cdot 2^{n^{k+k'}}\}$, the i -th interval is $[2(i-1) \cdot \frac{1}{2 \cdot 2^{n^{k+k'}}}, 2i \cdot \frac{1}{2 \cdot 2^{n^{k+k'}}}]$. M then makes $4 \cdot 2^{n^{k+k'}}$ oracle queries of the form (i, j) , where $i \in \{1, \dots, 2 \cdot 2^{n^{k+k'}}\}$ and $j \in \{0, 1\}$. Notice that the lengths of the queries are upper bounded by $2 + n^{k+k'}$, which is a polynomial as required by the class $\text{EXP}^{\text{poly}}\text{-NEXP}$.

For each query (i, j) , if $j = 0$ then the corresponding question is “whether there exists a strategy profile \vec{s} of the provers such that $u(\vec{s}; x)$ is in the i -th interval”; and if $j = 1$ then the corresponding question is “whether there exists a strategy profile \vec{s} such that $u(\vec{s}; x)$ is in the i -th interval and the first bit sent by P_1 is $c = 1$ ”. Notice that all the queries are indeed non-adaptive. We say that interval i is *non-empty* if the query $(i, 0)$ is answered 1, and *empty* otherwise.

Assume for now all the queries are answered correctly. M finds the highest index i^* such that the interval i^* is non-empty. It accepts if $(i^*, 1)$ is answered 1, and rejects otherwise. M clearly runs in exponential time. To see why it decides L (given correct oracle answers), notice that by Definition 1, there exists a strategy profile whose expected payment is non-negative and thus in $[0, 1]$ (since the payment is always in $[-1, 1]$). Thus there exists an interval i such that $(i, 0)$ is answered 1. Also by definition, the best strategy profile \vec{s} has the highest expected payment, and thus $u(\vec{s}; x)$ falls into interval i^* .

By Inequality 3, any strategy profile \vec{s}' with $u(\vec{s}'; x) < u(\vec{s}; x)$ has $u(\vec{s}'; x)$ not in interval i^* , since the difference between $u(\vec{s}'; x)$ and $u(\vec{s}; x)$ is larger than the length of the interval. Accordingly, all strategy profiles \vec{s}' with $u(\vec{s}'; x)$ in interval i^* satisfies $u(\vec{s}'; x) = u(\vec{s}; x)$: they are all the best strategy profiles of the provers. Again by Definition 1, P_1 sends the same first bit c under all these strategy profiles, $c = 1$ if and only if $x \in L$, and there does not exist any other strategy profile whose expected payment falls into interval i^* but the first bit sent by P_1 is different from c . Thus the answer to $(i^*, 1)$ always equals c , and M accepts if and only if $c = 1$, if and only if $x \in L$, as desired.

It remains to show that the oracle queries can be answered by an NEXP oracle. Recall that in the protocol (V, \vec{P}) a strategy s_{ij} of each prover P_i for each round j is a function mapping the transcript P_i has seen at the beginning of round j to the message he sends in that round. Since the protocol has polynomially many provers and polynomially many rounds, by definition a strategy profile consists of polynomially many functions from $\{0, 1\}^*$ to $\{0, 1\}^*$, where for each function both the input length and the output length are polynomial in n (since otherwise the messages cannot be processed by V). Accordingly, it takes exponentially many bits to specify each function: if the input length is $p(n)$ and the output length is $q(n)$, then $2^{p(n)}q(n)$ bits are sufficient to specify the truth table of the function. Therefore a strategy profile can also be specified by exponentially many bits. Below we construct an exponential-time non-deterministic

Turing machine M' that decides the questions corresponding to the queries.

Given any input $(i, 0)$, M' non-deterministically chooses a strategy profile \tilde{s} , in exponential time. It then goes through all the realizations of V 's random string, and for each realization r it runs V with x and r , and generates the provers' messages by looking them up in the corresponding truth table in \tilde{s} . M' computes for each r the payment output by V at the end of the protocol, and by combining these payments with corresponding probabilities M' computes the expected payment $u(\tilde{s}; x)$. If $u(\tilde{s}; x)$ is in interval i then M' accepts, otherwise it rejects. Since V 's random string is polynomially long, there are exponentially many realizations, and each one of them takes exponential time to run: V runs in polynomial time and it takes M' exponential time to look up the truth tables in \tilde{s} to generate a prover's message. Thus M' runs in non-deterministic exponential time. Also, if interval i is non-empty then there exists a strategy profile \tilde{s} that makes M' accept, and if interval i is empty then M' always rejects.

Similarly, given any input $(i, 1)$, M' non-deterministically chooses a strategy profile \tilde{s} and computes its expected payment $u(\tilde{s}; x)$. If $u(\tilde{s}; x)$ is not in interval i then M' rejects; otherwise, M' accepts if and only if the first bit sent by P_1 is 1. The correctness and the running time of this part follows immediately. \square

We have now established that $\text{MRIP} = \text{EXP}^{\|\text{poly}-\text{NEXP}\}$. To finish the proof of Theorem 1, we show that $\text{EXP}^{\|\text{poly}-\text{NEXP}\}$ equals $\text{EXP}^{\|\text{NP}\}$, by a padding argument. The detailed proof is omitted here and can be found in the full version [13].

LEMMA 8. $\text{EXP}^{\|\text{poly}-\text{NEXP}\} = \text{EXP}^{\|\text{NP}\}$.

PROOF OF THEOREM 1. Theorem 1 follows immediately from Lemmas 6, 7, and 8. \square

4. SIMULATING MRIP USING TWO PROVERS AND FIVE ROUNDS

In this section we prove Theorem 2, that is,

Theorem 2. $\text{MRIP} = \text{MRIP}(2, 5)$.

In particular, in Figure 5, we show how to simulate any MRIP protocol (V, \vec{P}) with $t(n)$ provers and $p(n)$ rounds using a verifier V' and two provers P'_1 and P'_2 . Recall that m_{ij} denotes the message exchanged between V and prover P_i in round j of their protocol. Essentially, V' asks one prover to simulate all other provers in the original protocol, and uses a second prover to cross-check his answers.

PROOF. Let (V, \vec{P}) be the MRIP protocol for a language L using $\vec{P} = (P_1, \dots, P_{t(n)})$ provers and $p(n)$ rounds of communication. Let \vec{m} denote the complete transcript of the protocol and m_{ij} be the message of prover P_i in round j of the protocol. We simulate this protocol using two provers P'_1 and P'_2 and verifier V' in the protocol (V', \vec{P}') in Figure 5. Let r denote the random coin flips used by V in protocol (V, \vec{P}) .

To see the correctness of the protocol, notice that,

- (a) Even though V' sends the random string r of V to P'_1 , he uses a different random string $r' \neq r$ to select a message in Step 4. Thus, the expected payment to the provers depends on r' , which is not known to them.
- (b) P'_2 does not know r and essentially commits to a transcript \vec{m}' of (V, \vec{P}) in Step 5.

The provers can lie in the following ways, with the corresponding payments:

- 1. P'_1 and P'_2 do not commit to the same transcript \vec{m} . Without loss of generality assume that, P'_1 lies on some message in \vec{m} . V catches this lie in Step 6a with probability $\frac{1}{p(n)t(n)}$ and gives a payment $R' = -1$.
- 2. P'_1 and P'_2 agree on the transcript \vec{m} , but \vec{m} does not correspond to the best strategy profile in protocol (V, \vec{P}) . In this case, they get a payment $\frac{R}{2p(n)t(n)}$, where R is the payment that the original protocol generates on this suboptimal transcript \vec{m} .

First of all, we can rule out Case 2 as a possible strategy of rational provers P'_1 and P'_2 . This is because V' pays them according to the original protocol (V, \vec{P}) . Committing to any dishonest transcript earns them a payment strictly less than the best possible. This follows from the correctness of the original MRIP protocol (V, \vec{P}) .

Now suppose P'_1 lies on $y > 0$ messages in \vec{m} —this corresponds to Case 1. The expected payment to P'_1 would then be

$$\begin{aligned} & -1 \left(\frac{y}{p(n)t(n)} \right) + \frac{R}{2p(n)t(n)} \left(\frac{p(n)t(n) - y}{p(n)t(n)} \right) \\ & \leq \frac{1}{p(n)t(n)} \left(\frac{R}{2} - y \right) < 0, \end{aligned}$$

where the last inequality is because $R \leq 1$ and $y \geq 1$. Thus, it is not rational for P'_1 to lie on any message in \vec{m} and the result follows. \square

REMARK 2. If we allow 3 provers instead of 2, then the verifier only needs to interact with each prover exactly once. In particular, V' sends r to P'_3 instead of P'_1 . This property may be desirable for applications that do not allow repeated interactions (e.g., each prover may be available only for a limited amount of time).

Optimal number of rounds.

We conclude this section with a discussion on the optimal number of rounds required to capture the full power of MRIP. It is well known that MIP can be simulated using two provers and one round [15], which is clearly optimal. It would be interesting to obtain a similar result for MRIP. Note that, because we count the number of rounds differently (provers' messages and verifier's queries are considered to be separate rounds) and we assume the first round of messages are always from the provers, any non-trivial MRIP protocol with at least two provers requires at least three rounds. Indeed, in less than three rounds, the verifier cannot cross-check answers with the second prover (beyond the initial messages received from them), in which case the protocol reduces to a single-prover protocol.

Currently, the protocol in Figure 5 requires two extra rounds because the verifier's queries to each prover are not parallel. The verifier uses the answers from its query to Prover 1 to form its query to Prover 2. Getting rid of this dependency may result in a 2-prover 3-round protocol for MRIP. Interestingly, note that our protocols for NEXP and coNEXP only require 3 rounds (Section 2.1), so we know $\text{NEXP} \cup \text{coNEXP} \subseteq \text{MRIP}(2, 3)$. We leave the following closely related problems open for future study: (1) what is the optimal number of rounds required to capture the full

Given any string x of length n ,

1. P'_1 sends c , which is the first bit of m_{11} , to V' . V' outputs c at the end of the protocol.
2. V' calculates the random bits r used by V in protocol (V, \vec{P}) . V' sends r to P'_1 .
3. P'_1 uses r to simulate the entire protocol (V, \vec{P}) and sends \vec{m} to V' .
4. V' chooses a round j at random from $\{1, \dots, p(n)\}$ and a prover index k from $\{1, \dots, t(n)\}$.
5. V' sends (j, k) and $\vec{m}_k = (m_{k1}, \dots, m_{k(j-1)})$ (that is, message transcript till $j - 1$ rounds of V with P_k) to P'_2 .
6. P'_2 simulates P_k on the round j and sends m'_{kj} to V .
7. V computes the payment R' as follows,
 - (a) If $j = 1$ and $k = 1$, then if the first bit of m'_{kj} is not c , set $R' = -1$.
 - (b) If $m_{kj} \neq m'_{kj}$, set $R' = -1$.
 - (c) Else, V' computes the payment R given by V in the protocol (V, \vec{P}) , using \vec{m} .
 V' then sets $R' = \frac{R}{2^{p(n)t(n)}}$.

Figure 5: Simulating MRIP with 2 provers and 5 rounds.

power of MRIP, and (2) what is the exact characterization of MRIP(2, 3)?

5. UTILITY GAPS IN MRIP

In our MRIP protocols until now, the provers are sensitive to arbitrarily small losses in the payment (similar to RIP in [5]). In [6] the authors strengthen their model by requiring the prover deviating from the honest protocol to suffer a non-negligible loss in their received payment (either constant or polynomial). This loss is demanded for *any* deviation from the prescribed behavior (i.e., the optimal strategy) and not just for reporting an incorrect answer to the membership of the input.

Formally, let \tilde{s} be the optimal strategy profile and \tilde{s}^* a suboptimal strategy profile of P . Then the *ideal* utility gap requires that $u(\tilde{s}) - u(\tilde{s}^*) > 1/\alpha(n)$, where $\alpha(n)$ is constant or polynomial in n . Although an ideal utility gap strongly guarantees that the prover uses his optimal strategy, as pointed out by [25], such a utility gap appears to be too strong to hold for many meaningful protocols.

To provide some intuition about the ideal utility gap, consider the MRIP protocol for NEXP given in Figure 2. This protocol does not satisfy the ideal utility gap condition, even though at first glance the gap appears to be significant. Indeed, the provers who *report the incorrect answer* always have a constant utility gap. However, a prover may tell the truth about the membership of x but deviate slightly from the optimal strategy. For example, the prover could lie so that the verifier accepts with probability $1 - \varepsilon$, where ε is exponentially small in $|x|$. In this situation, the verifier may have an exponentially small probability of detecting the deviation, and the expected payment would decrease by an exponentially small amount.

The difficulty in achieving the ideal gap can be seen in previous work as well. For example, the rational proof for MA in [6] fails to satisfy this constraint. Guo, Hubáček, Rosen and Vald in [25] also echo the concern about the ideal utility gap being strong and define a weaker notion of utility gap. However, they impose it on *rational arguments* rather than rational proofs, and they still consider a single prover. They also require that *noticeable deviation* leads to *noticeable loss*: if under a strategy \tilde{s}' of the prover, the probability for the verifier to output the correct answer is noticeably smaller than 1, then the expected payment to the prover under \tilde{s}' is also noticeably smaller than the optimal expected pay-

ment. While we do not impose this requirement, our notion of utility gap as defined below implies it, but not vice-versa.

In particular, our notion allows us to encompass both the protocol in [6] and the NEXP protocol in Figure 2. Intuitively, we require that the provers who report the membership of the input incorrectly suffer noticeable loss in their received payment.

DEFINITION 5 (UTILITY GAP). *Let L be a language in MRIP, (V, \vec{P}) an MRIP protocol for L , \tilde{s} the strategy profile of \vec{P} that maximizes the expected payment, and c the first bit sent by P_1 according to \tilde{s} . For any other strategy profile \tilde{s}' , let c' be the first bit sent by P_1 according to \tilde{s}' . We say that (V, \vec{P}) has an $\alpha(n)$ -utility gap if, whenever $c' \neq c$ we have*

$$\mathbb{E}_r R(x, r, (V, \vec{P})(x, r, \tilde{s})) - \mathbb{E}_r R(x, r, (V, \vec{P})(x, r, \tilde{s}')) > 1/\alpha(n).$$

We denote an MRIP protocol with an $\alpha(n)$ -utility gap as an $\alpha(n)$ -gap-MRIP protocol, and we shorten “polynomial in n ” as *poly*(n). It is easy to see that the protocol for NEXP in Figure 2 has an $O(1)$ -utility gap. Following the analysis of Lemma 1 we get the same for coNEXP. That is,

LEMMA 9. $\text{NEXP} \subseteq O(1)\text{-gap-MRIP}$ and $\text{coNEXP} \subseteq O(1)\text{-gap-MRIP}$.

We have shown that with an exponential utility gap, $\text{MRIP} = \text{EXP}^{\text{NP}}$. We now characterize MRIP protocols with a polynomial utility gap as exactly the class of languages decided by a polynomial-time oracle Turing machine with non-adaptive access to an NEXP oracle. That is,

Theorem 3. $\text{poly}(n)\text{-gap-MRIP} = \text{P}^{\text{NEXP}}$.

The proof of Theorem 3 uses similar ideas as that of Lemma 7 and can be found in the full version [13].

REMARK 3. *There is a tradeoff between the utility gap and the computational efficiency in the two MRIP protocols we constructed for NEXP: the protocol in Figure 2 has a constant utility gap, but it relies on the standard MIP protocol; the protocol in Figure 3 is very efficient, but it only has an exponential utility gap. It would be interesting to see if there exists an MRIP protocol for NEXP that has a noticeable utility gap and is (almost) as efficient as the protocol in Figure 3.*

6. REFERENCES

- [1] Amazon Mechanical Turk. Online at <https://www.mturk.com/mturk>.
- [2] Effective use of Amazon Mechanical Turk (MTurk). Online at <http://neerajkumar.org/writings/mturk/>.
- [3] Proof Market. Online at <https://proofmarket.org>.
- [4] S. Arora and B. Barak. *Computational complexity: a modern approach*. Cambridge University Press, 2009.
- [5] P. D. Azar and S. Micali. Rational proofs. In *Proceedings of the Forty-Fourth Annual Symposium on Theory of Computing (STOC)*, pages 1017–1028, 2012.
- [6] P. D. Azar and S. Micali. Super-efficient rational proofs. In *Proceedings of the Fourteenth Annual ACM Conference on Electronic Commerce (EC)*, pages 29–30, 2013.
- [7] L. Babai. Trading group theory for randomness. In *Proceedings of the Seventeenth Annual ACM Symposium on Theory of Computing (STOC)*, pages 421–429, 1985.
- [8] L. Babai, L. Fortnow, and C. Lund. Non-deterministic exponential time has two-prover interactive protocols. *Computational complexity*, 1(1):3–40, 1991.
- [9] L. Babai and S. Moran. Arthur-Merlin games: a randomized proof system, and a hierarchy of complexity classes. *Journal of Computer and System Sciences*, 36(2):254–276, 1988.
- [10] M. Ben-Or, S. Goldwasser, J. Kilian, and A. Wigderson. Multi-prover interactive proofs: How to remove intractability assumptions. In *Proceedings of the Twentieth Annual ACM Symposium on Theory of Computing (STOC)*, pages 113–131, 1988.
- [11] G. W. Brier. Verification of forecasts expressed in terms of probability. *Monthly weather review*, 78(1):1–3, 1950.
- [12] A. K. Chandra and L. J. Stockmeyer. Alternation. In *17th Annual Symposium on Foundations of Computer Science (FOCS)*, pages 98–108, 1976.
- [13] J. Chen, S. McCauley, and S. Singh. Rational proofs with multiple provers. *arXiv preprint arXiv:1504.08361*, 2015.
- [14] U. Feige and J. Kilian. Making games short. In *Proceedings of the Twenty-Ninth Annual ACM Symposium On Theory of Computing (STOC)*, pages 506–516, 1997.
- [15] U. Feige and L. Lovász. Two-prover one-round proof systems: their power and their problems. In *Proceedings of the Twenty-Fourth Annual ACM Symposium on Theory of Computing (STOC)*, pages 733–744, 1992.
- [16] U. Feige and A. Shamir. Multi-oracle interactive protocols with constant space verifiers. *Journal of Computer and System Sciences*, 44(2):259–271, 1992.
- [17] U. Feige, A. Shamir, and M. Tennenholtz. The noisy oracle problem. In *Proceedings of the Tenth Annual Conference on Advances in Cryptology (CRYPTO)*, pages 284–296, 1990.
- [18] J. Feigenbaum, D. Koller, and P. Shor. A game-theoretic classification of interactive complexity classes. In *Proceedings of Tenth Annual IEEE Structure in Complexity Theory Conference*, pages 227–237, 1995.
- [19] L. Fortnow, J. Rompel, and M. Sipser. On the power of multi-prover interactive protocols. *Theoretical Computer Science*, 134(2):545–557, 1994.
- [20] L. Fortnow and M. Sipser. Are there interactive protocols for co-NP languages? *Information Processing Letters (IPL)*, 28(5):249–251, 1988.
- [21] T. Gneiting and A. E. Raftery. Strictly proper scoring rules, prediction, and estimation. *Journal of the American Statistical Association*, 102(477):359–378, 2007.
- [22] O. Goldreich, S. Micali, and A. Wigderson. Proofs that yield nothing but their validity or all languages in NP have zero-knowledge proof systems. *Journal of the ACM (JACM)*, 38(3):690–728, 1991.
- [23] S. Goldwasser, S. Micali, and C. Rackoff. The knowledge complexity of interactive proof systems. *SIAM Journal on Computing*, 18(1):186–208, 1989.
- [24] S. Goldwasser and M. Sipser. Private coins versus public coins in interactive proof systems. In *Proceedings of the Eighteenth Annual ACM Symposium on Theory of Computing (STOC)*, pages 59–68, 1986.
- [25] S. Guo, P. Hubáček, A. Rosen, and M. Vald. Rational arguments: single round delegation with sublinear verification. In *Proceedings of the Fifth Annual Conference on Innovations in Theoretical Computer Science (ITCS)*, pages 523–540, 2014.
- [26] S. Guo, P. Hubáček, A. Rosen, and M. Vald. Rational sumchecks. In *13th Theory of Cryptography Conference (TCC'2016-A)*, to appear, 2016.
- [27] A. Kittur. Crowdsourcing, collaboration and creativity. *ACM Crossroads*, 17(2):22–26, 2010.
- [28] D. Koller and N. Megiddo. The complexity of two-person zero-sum games in extensive form. *Games and economic behavior*, 4(4):528–552, 1992.
- [29] C. Lund, L. Fortnow, H. Karloff, and N. Nisan. Algebraic methods for interactive proof systems. *Journal of the ACM (JACM)*, 39(4):859–868, 1992.
- [30] J. H. Reif. The complexity of two-player games of incomplete information. *Journal of Computer and System Sciences*, 29(2):274–301, 1984.
- [31] A. Shamir. $IP = PSPACE$. *J. ACM*, 39(4):869–877, 1992.
- [32] L. Von Ahn and L. Dabbish. Labeling images with a computer game. In *Proceedings of the SIGCHI conference on Human factors in computing systems*, pages 319–326, 2004.
- [33] L. Von Ahn and L. Dabbish. Designing games with a purpose. *Communications of the ACM*, 51(8):58–67, 2008.
- [34] R. Williams. Nonuniform ACC circuit lower bounds. *Journal of the ACM (JACM)*, 61(1):2, 2014.